
Special Topics in Security

ECE 5698

Engin Kirda
ek@ccs.neu.edu



Northeastern University

Low-Level Essentials for Understanding Security Problems

Computer Architecture

- The modern computer architecture is based on “Von Neuman”
 - Two main parts: CPU (Central Processing Unit) and Memory
 - This architecture is used everywhere (e.g., even mobile phones)
 - This architecture is fundamental, has not changed yet
- What is memory used for?
 - E.g., location of curser, size of windows, shape of each letter being displayed, graphics of icons, text, values, etc...
 - Von Neuman also says that not only data, but programs should also be in main memory

The CPU

- Storing data by itself, of course, is not enough
 - The CPU reads instructions from memory one by one
 - Executes them (*fetch-execute-cycle*)
- Following components make up the CPU
 - Program Counter
 - Instruction Decoder
 - Data bus
 - General-purpose registers
 - Arithmetic and logic unit

Program Counter and Instruction Decoder

- Is used to tell the CPU where to fetch next instruction
 - There is no difference between memory and data
 - Program counter holds memory address of next instruction
- The instruction decoder then makes sense of the instruction
 - Addition? Subtraction? Multiplication? Move operation? Etc.
 - A typical instruction usually consists of memory locations as well
 - E.g., move this piece of data from memory address X to memory address Y

The Data Bus and Registers

- The data bus connects the memory and the CPU
 - i.e., it is the actual, physical wire
- In addition to the memory, the processor has high-speed, special memory location
 - Called *registers*
 - Special-purpose registers
 - General-purpose registers
 - Registers are used for computation

Arithmetic and Logic Unit

- Once instruction has been decoded, CPU passes data and decoded instruction to ALU
 - Now, the instruction is actually executed
 - Results of the computation are placed on the data bus and sent to memory locations given in the instruction
 - For example, we can tell ALU to add 1 to register A and place result in register B

Some basics (you know them ;))

- The number attached to each storage location
 - ... is an *address*
 - A single storage location is called a *byte*
 - On x86 processors, a byte is between 0..255
 - Obviously, two bytes can be used to represent any number between 0..65536
 - Four bytes can be used to represent numbers between 0..4294967295., Luckily, we do not have to worry about this. The architecture helps us to do math with 4 byte numbers
- Addresses sizes are in *words* (a word is 4 bytes)
 - Note that this means that a number and an address are dealt with the same way

Data Accessing Methods

- Processors have a number of different ways of accessing data
 - Known as *addressing modes*
 - The simplest method is known as *immediate mode*
 - Data access is enabled in instruction itself
 - In *register addressing mode* the instruction contains a register to access rather than memory location
 - In *direct addressing mode*, the instruction contains memory address
 - In *indexed addressing mode*, the instruction contains memory address to access and an index register to offset
 - E.g., Address 1000, register=4, address=1004

Data Accessing Methods

- Processors have a number of different ways of accessing data
 - Indirect addressing mode
 - We take address that is stored in register
 - Base point addressing mode
 - You take address in register and add an offset
 - Used a lot

A Simple Program in Assembler

```
# No INPUT
# Returns a status code, you can view it by typing echo $?
# %ebx holds the return code

.section .data
.section .text
.globl _start

_start:
mov $1, %eax # This is the sys call for exiting program
movl $0, %ebx # This value is returned as status
int $0x80 # This interrupt calls the kernel, to execute sys call
```

So now we have the source code...

- So how do we create the application?
 - Well, we need to assemble and link the code
 - This can be done by using the assembler *as*:
 - *as exit.s -o exit.o*
 - *ld -o exit exit.o*

Anatomy of the Code

- `.section` breaks up the program into sections
 - `.data`, obviously, is used for variables and data you might need
 - `.text` is the code of the program that you write
 - `.globl __start` indicates that `__start` is a symbol, required by the linker
 - `__start`, in fact, indicates that the loader will load and start the program from this location
 - The next instruction, *movl*, has two operands: *source* and *destination*
 - *movl, addl, subl*

Registers on the x86 Architecture

- General purpose registers:
 - `%eax`, `%ebx`, `%ecx`, `%edx`, `%edi`, `%esi`
- Special purpose registers:
 - `%ebp`, `%esp`, `%eip`, `%eflags`
- Some registers (e.g., `%eip`, `%eflags`) can only be accessed through special instructions
- The \$ sign before a number means that we are using *immediate addressing mode*

Let us write a more complicated Program

- Task: Find the maximum of a list of numbers
 - Questions to ask:
 - Where will the numbers be stored?
 - How do we find the maximum number?
 - How much storage do we need?
 - Will registers be enough or is memory needed?
 - Let us designate registers for the task at hand:
 - %edi holds position in list
 - %ebx will hold current highest
 - %eax will hold current element examined

“Algorithm” we use

- Check if %eax is zero (i.e., termination sign)
 - If yes, exit
 - If not, increase current position %edi
 - Load next value in the list to %eax
 - We need to think about what addressing mode to use here
 - Compare %eax (the current value) with highest value so far %ebx
 - If the current value is higher, replace %ebx
 - Repeat until termination

Let's get down to the code

```
.section .data
    data_items:
        .long 3,67,34,222,45,75,54,34,44,33,22,11,66,0

.section .text
.globl _start
_start:
    mov $0, %edi # Reset index
    movl data_items(,%edi,4), %eax
    movl %eax, %ebx #First item is the biggest so far
start_loop:
    cmpl $0, %eax
    je loop_exit
```

Let's get down to the code

```
incl %edi # Increment edi
movl data_items(,%edi,4), %eax #load the next value
cmpl %ebx, %eax # Compare ebx with eax
jle start_loop # if it is less, then just jump to the beginning
movl %eax, %ebx #Otherwise, store the new largest number
jmp start_loop


loop_exit:
movl $1, %eax # Remember the exit sys call? It is 1
int $0x80
```

Important Instruction

- The compare instruction
 - `cmpl $0, %eax`
 - `je end_loop`
 - Other jump instructions *`jg, jge, jl, jle, jmp`*
- **mov** instruction
 - Used often. One of the most important and common instructions that you are going to see
 - ... and use, for example, when writing shell code

What if we do not want to move a word?

- Suppose you only need to move a *byte* at a time, not a word
 - You can use the *movb* instruction
 - In %eax, the least significant half is addressed by %ax (i.e., 2 bytes)
 - %ax is divided into %ah and %al (1 byte in size)



Functions

- A function is composed of several different pieces
 - function name
 - Symbol that represents where the function starts
 - function parameters
 - Data items passed to function for processing
 - Local variables
 - Temporary storage areas used in the function
 - Thrown away when the processing finishes
 - Static variables
 - Storage area that is reused over invocations
 - Global variables
 - Storage areas outside the function

The Return Address

- The return address is a parameter that tells the function where to resume executing after the function is completed
 - It is “invisible” – the programmer does not necessarily know the address
 - When a function is invoked, the calling point is saved
 - When the function completes, it returns to the initial calling point
 - In machine code, functions are called with *call* and they return when they execute *ret*
- Return value
 - Usually, a single value is returned to caller

The Calling Convention

- The way that the variables are stored and the parameters and return values are transferred by the computer varies from language to language as well
 - This variance is known as a language's *calling convention*
- The assembly language can use any calling convention
 - You can make one up yourself... however...
 - If you want to interoperate with other languages, you need to follow their calling conventions
 - e.g., suppose you want your code to be callable from a C program...

The Stack

- Each computer program that runs uses a region of memory called the stack to enable functions to work properly
 - You generally keep the things that you are working on toward the top, and you take things off as you are finished working with them
- The computer's stack lives at the very top addresses of memory
 - You can push values onto stack using *pushl*
 - You can retrieve items from the stack using *popl*

The Stack

- Where is the “top” of the stack
 - Because of architectural considerations, the computer stack grows from higher addresses to lower addresses
 - i.e., it grows downwards
 - How do we know where the “top” of the stack is?
 - The `%esp` register stores a pointer to stack location
 - If something is pushed onto stack, the stack decreases by `%esp – 4` if *pushl* is used
 - Yes, but if I only want to read? How can I do this without popping?
 - `movl (%esp), %eax`
 - `movl 4(%esp), %eax` (for addressing a higher value)

The C calling convention

- The stack is the key element for implementing a function's...
 - local variables, parameters, and return address
 - Before executing a function, a program pushes all of the parameters for the function onto the stack in the reverse order that they are documented
 - Program issues *call* instruction
 - *call* does two things
 - pushes address of next instruction (i.e., return)
 - Modifies %eip to point to function start

The C calling convention

- Parameter #N

...

Parameter 2

Parameter 1

Return Address <--- (%esp)

- Now, function has to do some thing
 - It saves the current base pointer register %ebp
 - Needed for local variables and parameters
 - `pushl %ebp`
 - `movl %esp, %ebp`

The C calling convention

- Copying the stack pointer into the base pointer at the beginning of a function allows you to always know where your parameters are

Parameter #N <--- $N*4+4(\%ebp)$

...

Parameter 2 <--- $12(\%ebp)$

Parameter 1 <--- $8(\%ebp)$

Return Address <--- $4(\%ebp)$

Old %ebp <--- $(\%esp)$ and $(\%ebp)$

The C calling convention

- *stack frame* consists of all of the stack variables used within a function
- Next, the function reserves space on the stack for any local variables it needs
 - This is done by simply moving the stack pointer out of the way
 - E.g., if we need two words: *subl \$8, %esp*
 - Variables are local because when function returns, the stack frame is reset and variables disappear

The C calling convention

- Our stack now looks like this:
 - Parameter #N <--- $N*4+4(\%ebp)$
 - ...
 - Parameter 2 <--- $12(\%ebp)$
 - Parameter 1 <--- $8(\%ebp)$
 - Return Address <--- $4(\%ebp)$
 - Old %ebp <--- $(\%ebp)$
 - Local Variable 1 <--- $-4(\%ebp)$
 - Local Variable 2 <--- $-8(\%ebp)$ and $(\%esp)$

The C calling convention

- Now we can use the `%ebp` for accessing all variables
 - `%ebp` was specifically designed for this purpose
- When a function is done...
 - It stores its return value in `%eax`
 - It resets the stack to what it was when it was called
 - It gets rid of the current stack frame and puts the stack frame of the calling code back into effect
 - It returns by invoking the `ret` command
 - `movl %ebp, %esp`
`popl %ebp`
`ret`

Buffer Overflows

Buffer Overflows

- A buffer overflow occurs any time the program attempts to store data beyond the boundaries of a buffer, overwriting the adjacent memory locations
 - Results from mistakes while writing the code
 - Unfamiliarity with the language
 - Ignorance about security issues
 - Unwillingness to take any extra effort
- Vulnerable software
 - Mostly C/C++ programs
 - Not language with automated memory management
 - Java, Perl, Python, C#

Buffer Overflows

- Common targets
 - setuid/setgid programs
 - network servers
- Goals
 - Overwrite other “interesting” variables
(file names, passwords, pointers...)
 - Force the program to execute operations it was not intended to do:
 - inject (or simply find) code into the process memory
 - change flow of control (flow of execution) to execute that code

Buffer Overflows


- Morris worm (1988): overflow in fingerd
 - 6,000 machines infected (10% of the Internet)
- CodeRed (2001): overflow in MS-IIS server
 - 300,000 machines infected in 14 hours
- SQL Slammer (2003): overflow in MS-SQL server
 - 75,000 machines infected in **10 minutes**
- In 2003, around 75% of the vulnerabilities were buffer overflows (CERT)
 - Why are they still interesting and important?

Part I

A deeper look into the STACK

Memory Layout

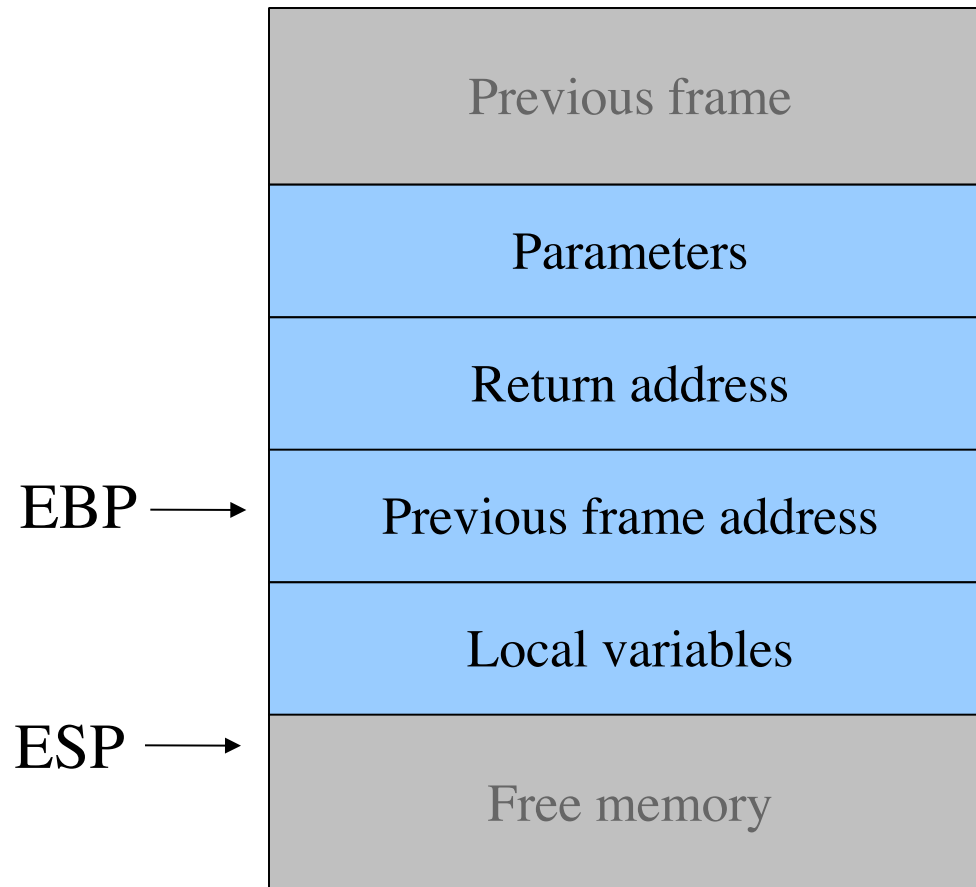
- Stack segment
 - local variables
 - procedure activation records
- Data segment
 - global uninitialized variables (.bss)
 - global initialized variables (.data)
 - dynamic variables (heap)
- Code (Text) segment
 - program instructions
 - usually read-only
- In linux, under the proc filesystem
 - `>cat /proc/<pid>/maps`



The Stack

- In most architectures (Intel, Motorola, Sparc), stack grows towards bottom
- The ESP (stack pointer) register points to the top of the stack
- The EBP (base pointer) points to the current frame
- Each frame contains:
 - Return address (where to jump at the end of the function)
 - Address of the previous frame
 - Parameters
 - Local variables of the function

Frame




Procedure Call

```
int foo(int a, int b)
{
    int i = 3;

    return (a + b) * i;
}
```

```
int main()
{
    int e = 0;
    e = foo(4, 5);
    printf("%d", e);
}
```




A closer look

```
(gdb) disas main
```

```
Dump of assembler code for function main:
```

```
0x0804836d <main+0>:    push    %ebp
0x0804836e <main+1>:    mov     %esp,%ebp
0x08048370 <main+3>:    sub     $0x18,%esp
0x08048373 <main+6>:    and     $0xfffffffff0,%esp
0x08048376 <main+9>:    mov     $0x0,%eax
0x0804837b <main+14>:   add     $0xf,%eax
0x0804837e <main+17>:   add     $0xf,%eax
0x08048381 <main+20>:   shr     $0x4,%eax
0x08048384 <main+23>:   shl     $0x4,%eax
0x08048387 <main+26>:   sub     %eax,%esp
0x08048389 <main+28>:   movl    $0x0,0xffffffffc(%ebp)
0x08048390 <main+35>:   movl    $0x5,0x4(%esp)
0x08048398 <main+43>:   movl    $0x4,(%esp)
0x0804839f <main+50>:   call    0x8048354 <foo>
0x080483a4 <main+55>:   mov     %eax,0xffffffffc(%ebp)
```



A closer look

```
(gdb) breakpoint foo
Breakpoint 1 at 0x804835a
(gdb) run
Starting program: ./test1
Breakpoint 1, 0x0804835a in foo ()
(gdb) disas
Dump of assembler code for function foo:
0x08048354 <foo+0>:      push    %ebp
0x08048355 <foo+1>:      mov     %esp,%ebp
0x08048357 <foo+3>:      sub     $0x10,%esp
0x0804835a <foo+6>:      movl    $0x3,0xffffffffc(%ebp)
0x08048361 <foo+13>:     mov     0xc(%ebp),%eax
0x08048364 <foo+16>:     add     0x8(%ebp),%eax
0x08048367 <foo+19>:     imul    0xffffffffc(%ebp),%eax
0x0804836b <foo+23>:     leave
0x0804836c <foo+24>:     ret
End of assembler dump.
(gdb)
```

5
4
0x080483a4
0xafdde9f8
3

The “foo” frame

```
(gdb) stepi
0x08048361 in foo ()
(gdb) x/12wx $ebp-16
0xaf9d3cc8: 0xaf9d3cd8 0x080482de 0xa7faf360 0x00000003
0xaf9d3cd8: 0xafdde9f8 0x080483a4 0x00000004 0x00000005
0xaf9d3ce8: 0xaf9d3d08 0x080483df 0xa7fadf14 0x08048430
```

