
Special Topics in Security

ECE 5698

Engin Kirda
ek@ccs.neu.edu



Northeastern University

Race Conditions

Overview

- Parallel execution of tasks
 - multi-process or multi-threaded environment
 - tasks can interact with each other
- Interaction
 - shared memory (or address space)
 - file system
 - signals
- Results of tasks depends on relative timing of events
→ Indeterministic behavior

Race Conditions

- Race conditions
 - alternative term for indeterministic behavior
 - often a robustness issue
 - but also many important security implications
- Assumption needs to hold for some time for correct behavior, but assumption can be violated
- Time window when assumption can be violated
 - window of vulnerability

Race Conditions

- Window of vulnerability can be very short
 - race condition problems are difficult to find with testing and difficult to reproduce
 - attacker can slow down victim process/machine to extend window and can often launch many attempts
- Deadlock
 - special form of race condition
 - two processes are preventing each other from accessing a shared resource, resulting in both processes ceasing to function

Race Conditions

- General assumption
 - sequence of operations
 - is not atomic
 - can be interrupted at any time for arbitrary lengths
 - use proper countermeasures to ensure deterministic results

→ Synchronization primitives
- Locking
 - can impose performance penalty
 - critical section has to be as small as possible

Race Conditions

- Case study

```
public class Counter extends HttpServlet {  
    int count = 0;  
    public void doGet(HttpServletRequest in,  
                      HttpServletResponse out)  
    {  
        out.setContentType("text/plain");  
        PrintWriter p = out.getWriter();  
        count++;  
        p.println(count + " hits so far!");  
    }  
}
```

Race Conditions

- Time-of-Check, Time-of-Use (TOCTOU)
 - common race condition problem
 - problem:
 - Time-Of-Check** (t_1): validity of assumption A on entity E is checked
 - Time-Of-Use** (t_2): assuming A is still valid, E is used
 - Time-Of-Attack** (t_3): assumption A is invalidated
- Program has to execute with elevated privilege
 - otherwise, attacker races for his own privileges

TOCTOU

- Steps to access a resource
 1. obtain reference to resource
 2. query resource to obtain characteristics
 3. analyze query results
 4. if resource is fit, access it
- Often occurs in Unix file system accesses
 - check permissions for a certain file name (e.g., using **access(2)**)
 - open the file, using the file name (e.g., using **fopen(3)**)
 - four levels of indirection (symbolic link - hard link - inode - file descriptor)
- Windows uses file handles and includes checks in API open call

Overview

- Case study

```
/* access returns 0 on success */  
if(!access(file, W_OK)) {  
    ...  
    f = fopen(file, "wb+");  
    write_to_file(f);  
    ...  
} else {  
    fprintf(stderr, "Permission denied when trying to open %s.\n", file);  
}
```

- Attack

```
$ touch dummy; ln -s dummy pointer  
$ rm pointer; ln -s /etc/passwd pointer
```

Examples

- TOCTOU Examples
 - Filename Redirection
 - Setuid Scripts
 1. `exec()` system call invokes `setuid()` call prior to executing program
 2. program is a script, so command interpreter is loaded first
 3. program interpreted (with root privileges) is invoked on script name
 4. attacker can replace script content between step 2 and 3

Examples

- “Vulnerability” in certain browsers (Firefox, Opera)
 - user is registering for something and is asked to type “ONLY”
 - when “L” is pressed, security relevant application is started, user then presses “Y”...
- User is tricked into double-clicking a certain area in the browser
 - What happens? User is tricked into clicking on “Y” and a malicious application (plug-in) is installed...
- Such vulnerabilities are difficult to discover and fix
 - one way of fixing could be to build in delays and randomly place dialogs on the screen

Examples

- TOCTOU Examples
 - Directory operations
 - **rm** can remove directory trees, traverses directories depth-first
 - issues **chdir("../")** to go one level up after removing a directory branch
 - by relocating subdirectory to another directory, arbitrary files can be deleted
 - SQL **select** before **insert**
 - when select returns no results, insert a (unique) element
 - when DB does not check, possible to insert two elements with same key

Examples

- TOCTOU Examples
 - LOMAC
 - Linux kernel level monitor
 - checks system calls
 - arguments copied to module and checked
 - then, arguments are copied again to invoke actual system call

Examples

- TOCTOU Examples
 - File meta-information
 - **chown(2)** and **chmod(2)** are unsafe
 - operate on file names
 - use **fchown(2)** and **fchmod(2)** that use file descriptors
 - Joe Editor
 - when joe crashes (e.g., segmentation fault, xterm crashes)
 - unconditionally append open buffers to local DEADJOE file
 - DEADJOE could be symbolic link to security-relevant file

Temporary Files

- Similar issues as with regular files
 - commonly opened in **/tmp** or **/var/tmp**
 - often guessable file names
- “Secure” procedure
 1. pick a prefix for your filename
 2. generate at least 64 bits of high-quality randomness
 3. base64 encode the random bits
 4. concatenate the prefix with the encoded random data
 5. set umask appropriately (0066 is usually good)
 6. use **fopen(3)** to create the file, opening it in the proper mode
 7. delete the file immediately using **unlink(2)**
 8. perform reads, writes, and seeks on the file as necessary
 9. finally, close the file

Temporary Files

- Library functions to create temporary files can be insecure
 - **mktemp(3)** is not secure, use **mkstemp(3)** instead
 - old versions of **mkstemp(3)** did not set umask correctly
- Temp Cleaners
 - programs that clean “old” temporary files from **temp** directories
 - first **lstat(2)** file, then use **unlink(2)** to remove files
 - vulnerable to race condition when attacker replaces file between **lstat(2)** and **unlink(2)**
 - arbitrary files can be removed
 - delay program long enough until temp cleaner removes active file

Prevention

- “Handbook of Information Security Management” suggests
 1. increase number of checks
 2. move checks closer to point of use
 3. immutable bindings
- Only number 3 is acceptable!
- Immutable bindings
 - operate on file descriptors
 - do not check access by yourself (i.e., no use of **access(2)**)
drop privileges instead and let the file system do the job
- Use the **O_CREAT | O_EXCL** flags to create a new file with **open(2)** and be prepared to have the open call fail

Prevention

- Some calls require file names
`link()`, `mkdir()`, `mknod()`, `rmdir()`, `symlink()`, `unlink()`
 - especially **`unlink(2)`** is troublesome
- Secure File Access
 - create “secure” directory
 - directory only write and executable by UID of process
 - check that no parent directory can be modified by attacker
 - walk up directory tree checking for permissions and links at each step

Locking

- Ensures exclusive access to a certain resource
- Used to circumvent accidental race conditions
 - advisory locking (processes need to cooperate)
 - not mandatory, therefore not secure
- Often, files are used for locking
 - portable (files can be created nearly everywhere)
 - “stuck” locks can be easily removed
- Simple method
 - open file using the `O_EXCL` flag

Non-FS Race Conditions

- Linux / BSD kernel ptrace(2) / execve(2) race condition
- ptrace(2)
 - debugging facility
 - used to access other process' registers and memory address space
 - can only attach to processes of same UID, except being run by root
- execve(2)
 - execute program image

Non-FS Race Conditions

- Problem with **execve(2)**
 1. first checks whether process is being traced
 2. open image (may block)
 3. allocate memory (may block)
 4. set process EUID according to setuid flags
- Window of vulnerability between step 1 and step 4
 - attacker can attach via ptrace
 - blocking kernel operations allow other user processes to run

Non-FS Race Conditions

- Signaler handler race conditions
- Signals
 - used for asynchronous communication between processes
 - signal handler can be called in response to multiple signals
 - signal handler must be written re-entrant or block other signals
- Example
 - sendmail up to 8.11.3 and 8.12.0.Beta7
 - syslog(3) is called inside the signal handler
 - race condition can cause heap corruption because of double free vulnerability

Non-FS Race Conditions

- Windows DCOM / RPC vulnerability
 - RPCSS service
 - multiple threads process single packet
 - one thread frees memory,
while other process still works on it
 - can result in memory corruption
 - and thus denial of service

Detection

- Static code analysis
 1. specify potentially unsafe patterns
and perform pattern matching on source code
 2. source code analysis and model checking
 - MOPS (MOdel-checking Programs for Security properties)

Detection

- Static code analysis
 3. Source code analysis and annotations / rules
 - RacerX (found problems in Linux and commercial software)
 - rccjava (found problems in java.io and java.util)
- Dynamic analysis
 1. inferring data races during runtime
 - “Eraser: A Dynamic Data Race Detector for Multithreaded Programs” , ACM Transactions on CS, 1997

...and to complete...
Testing

Overview

- When system is designed and implemented
 - correctness has to be **tested**
- Different types of tests are necessary
 - validation
 - is the system designed correctly?
 - does the design meet the problem requirements?
 - verification
 - is the system implemented correctly?
 - does the implementation meet the design requirements?
- Different features can be tested
 - functionality, performance, *security*

Testing

- Edsger Dijkstra

Program testing can be quite effective for showing the presence of bugs, but is hopelessly inadequate for showing their absence.

- Testing
 - analysis that discovers what *is* and compares it to what *should be*
 - should be done throughout the development cycle
 - necessary process
 - but not a substitute for sound design and implementation
 - for example, running public attack tools against a server cannot prove that server is implemented securely

Testing

- Classification of testing techniques
 - white-box testing
 - testing all the implementation
 - path coverage considerations
 - faults of commission
 - find implementation flaws
 - but cannot guarantee that specifications are fulfilled
 - black-box testing
 - testing against specification
 - only concerned with input and output
 - faults of omissions
 - specification flaws are detected
 - but cannot guarantee that implementation is correct

Testing

- Classification of testing techniques
 - static testing
 - check requirements and design documents
 - perform source code auditing
 - theoretically reason about (program) properties
 - cover a possible infinite amount of input (e.g., use ranges)
 - no actual code is executed
 - dynamic testing
 - feed program with input and observe behavior
 - check a certain number of input and output values
 - code is executed (and must be available)

Testing

- Automatic testing
 - testing should be done continuously
 - involves a lot of input, output comparisons, and test runs
 - therefore, ideally suitable for automation
 - testing hooks are required, at least at module level
 - nightly builds with tests for complete system are advantageous
- Regression tests
 - test designed to check that a program has not "regressed", that is, that previous capabilities have not been compromised by introducing new ones

Testing

- Software fault injection
 - go after effects of bugs instead of bugs
 - reason is that bugs cannot be completely removed
 - thus, make program fault-tolerant
 - failures are deliberately injected into code
 - effects are observed and program is made more robust
- Most testing techniques can be used to identify security problems

Security Testing

- Design level
 - not much tool support available
 - manual design reviews
 - formal methods
 - attack graphs
- Formal methods
 - formal specification that can be mathematically described and verified
 - often used for small, *safety*-critical programs
e.g., control program for nuclear power plant
 - state and state transitions must be formalized and unsafe states must be described
 - “model checker” can ensure that no unsafe state is reached

Security Testing

- Attack graph
 - given
 - a finite state model, M , of a network
 - a security property P
 - an attack is an execution of M that violates P
 - an attack graph is a set of attacks of M
- Attack graph generation
 - done by hand
 - error prone and tedious
 - impractical for large systems
 - automatic generation
 - provide state description
 - transition rules


Sandia Red Team “White Board” attack graph
from DARPA CC20008 Information battle space
preparation experiment



Security Testing

P = Attacker gains root access to Host 1.

4 hosts
30 actions
310 nodes
3400 edges



Security Testing

- Implementation Level
 - detect known set of problems and security bugs
 - more automatic tool support available
 - target particular flaws
 - reviewing (auditing) software for flaws is reasonably well-known and well-documented
 - support for static and dynamic analysis
 - ranges from “how-to” for manual code reviewing to elaborate model checkers or compiler extensions

Static Security Testing

- Manual auditing
 - code has to support auditing
 - architectural overview
 - comments
 - functional summary for each method
 - OpenBSD is well know for good auditing process
 - 6 -12 members since 1996
 - comprehensive file-by-file analysis
 - multiple reviews by different people
 - search for bugs in general
 - proactive fixes
 - Microsoft also has intensive auditing processes
 - Every piece of written code has to be reviewed by another developer

Static Security Testing

- Manual auditing
 - tedious and difficult task
 - some initiatives were less successful
 - Sardonix (security portal)
*“Reviewing old code is tedious and boring and no one wants to do it,”
Crispin Cowan said.*
 - Linux Security Audit Project (LSAP)

Statistics for All Time

Lifespan		Rank		Page Views		D/1		Bugs		Support		Patches		Trkr		Tasks
1459 days		0 (0.00)		4,887		0		0 (0)		0 (0)		0 (0)		0 (0)		0 (0)

Static Security Testing

- Syntax checker
 - parse source code and check for functions that have known vulnerabilities, e.g., `strcpy()`, `strcat()` (as we saw in the buffer overflows lecture)
 - also limited support for arguments (e.g., variable, static string)
 - only suitable as first basic check
 - cannot understand more complex relationships
 - no control flow or data flow analysis
- Examples
 - flawfinder
 - RATS (Rough Auditing Tool for Security)
 - ITS4

Static Security Testing

- Annotation-based systems
 - programmer uses annotations to specify properties in the source code (e.g., this value must not be NULL)
 - analysis tool checks source code to find possible violations
 - control flow and data flow analysis is performed
 - Examples
 - SPlint
 - Eau-claire
 - UNO (uninitialized vars, out-of-bounds access)

Static Security Testing

- Model-checking
 - programmer specifies security properties that have to hold
 - models realized as state machines
 - statements in the program result in state transitions
 - certain states are considered insecure
 - usually, control flow and data flow analysis is performed
 - example properties
 - drop privileges properly
 - race conditions
 - creating a secure chroot jail
 - examples
 - MOPS (an infrastructure for examining security properties of software)